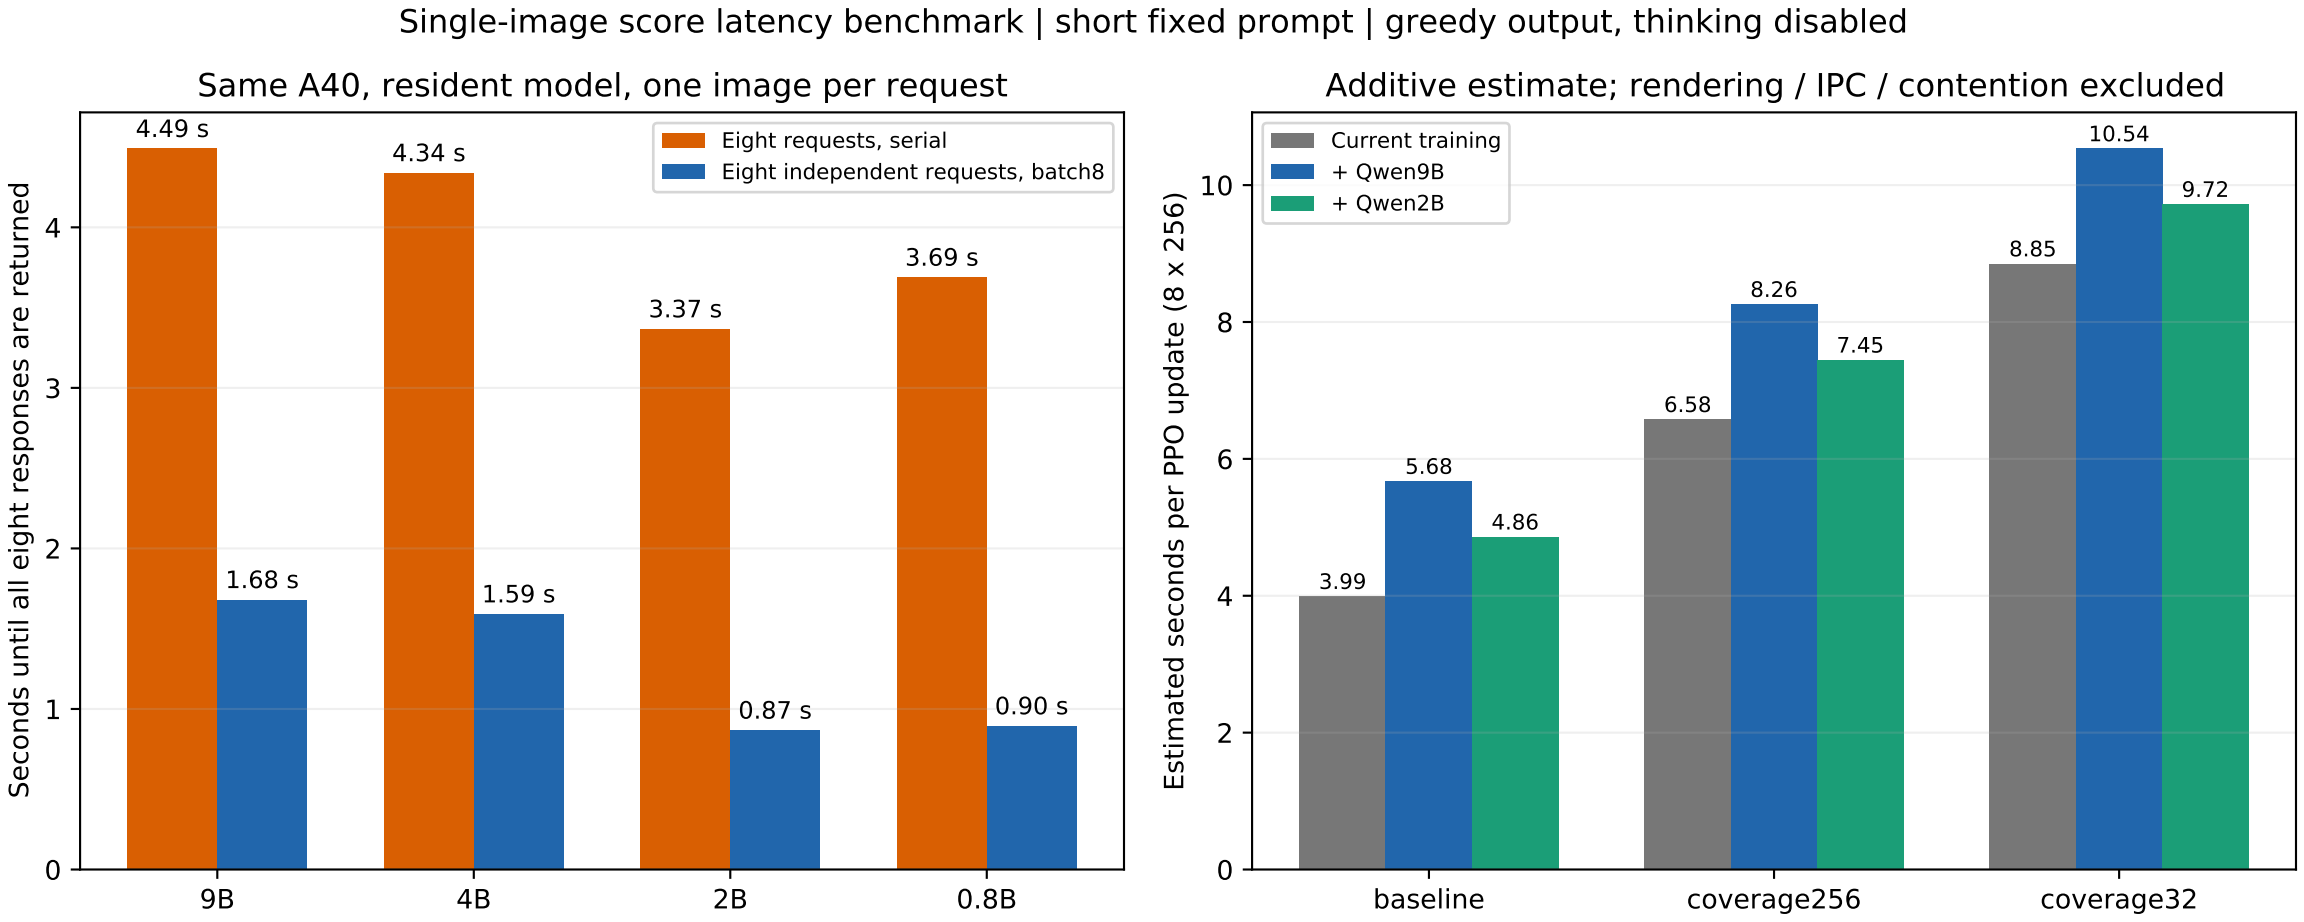
Single-image score latency benchmark | short fixed prompt | greedy output, thinking disabled
Same A40, resident model, one image per request
Seconds until all eight responses are returned
Eight requests, serial
Eight independent requests, batch8
4.49 s
1.68 s
4.34 s
1.59 s
3.37 s
0.87 s
3.69 s
0.90 s
9B
4B
2B
0.8B
Additive estimate; rendering / IPC / contention excluded
Estimated seconds per PPO update (8 x 256)
Current training
+ Qwen9B
+ Qwen2B
3.99
5.68
4.86
6.58
8.26
7.45
8.85
10.54
9.72
baseline
coverage256
coverage32